

CLOUD COMPUTING
– CSC8110
Coursework 2023-24
Project Report

Ronil Rodrigues
MSc Cloud Computing
230112209

Table of Contents

AIM	3
IMPLEMENTATION	3
CONCLUSION.....	12
REFERENCES	12

AIM:

This coursework helps you to learn and understand the fundamentals of programming and deploying Kubernetes-based (an emerging cloud virtualisation technology) application hosting environment. By successfully completing the coursework, you will be able to gain hands-on experience in the following interrelated aspects including:

- Configuring a Kubernetes-based application hosting environment;
- Building, pushing and pulling images from the Docker Hub (global repository of software components' images maintained by the developers);
- Creating and deploying a complex web application stack consisting of multiple software components (e.g., web server, monitoring, etc.);

IMPLEMENTATION:

TASK 1: DEPLOY AND ACCESS THE KUBERNETES DASHBOARD AND A WEB APPLICATION COMPONENT

A sample Java web application component image, name "nclcloudcomputing/javabenchmarkapp ", has been uploaded to the Docker Hub. The image contains a ready-to-use implementation of a web application deployed in a Tomcat server (an open-source web server). In terms of computational logic, the web application implements a prime number check on a large number. By doing so, the application can generate high CPU and memory load.

1. Deploy 'Kubernetes Dashboard' on the provided VM with CLI and access/login the Dashboard.

We deploy Kubernetes Dashboard using the below command.

```
student@edge:~$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.7.0/aio/deploy/recommended.yaml
namespace/kubernetes-dashboard unchanged
serviceaccount/kubernetes-dashboard unchanged
service/kubernetes-dashboard unchanged
secret/kubernetes-dashboard-certs unchanged
secret/kubernetes-dashboard-csrf configured
Warning: resource secrets/kubernetes-dashboard-key-holder is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply.
Resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.
secret/kubernetes-dashboard-key-holder configured
configmap/kubernetes-dashboard-settings unchanged
role.rbac.authorization.k8s.io/kubernetes-dashboard unchanged
clusterrole.rbac.authorization.k8s.io/kubernetes-dashboard unchanged
rolebinding.rbac.authorization.k8s.io/kubernetes-dashboard unchanged
clusterrolebinding.rbac.authorization.k8s.io/kubernetes-dashboard unchanged
deployment.apps/kubernetes-dashboard unchanged
service/dashboard-metrics-scraper unchanged
deployment.apps/dashboard-metrics-scraper unchanged
```

Once the Dashboard is deployed, we need to follow the below steps to gain access to the Kubernetes dashboard.

```
dashboard.yaml
1 apiVersion: v1
2 kind: ServiceAccount
3 metadata:
4   name: admin-user
5   namespace: kubernetes-dashboard
6
```

In the above file we define the service account with the name admin-user in namespace Kubernetes-dashboard, we execute the service using “kubectl apply -f Kubernetes-dashboard.yaml” command.

```
ClusterRoleBinding.yaml
1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: ClusterRoleBinding
3 metadata:
4   name: admin-user
5 roleRef:
6   apiGroup: rbac.authorization.k8s.io
7   kind: ClusterRole
8   name: cluster-admin
9 subjects:
10 - kind: ServiceAccount
11   name: admin-user
12   namespace: kubernetes-dashboard
13
```

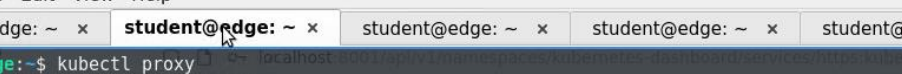
ClusterRoleBinding binds the cluster-admin ClusterRole to the admin-user ServiceAccount in the kubernetes-dashboard namespace. This configuration grants superuser access within the kubernetes-dashboard namespace to the admin-user ServiceAccount. We execute the service using “kubectl apply -f ClusterRoleBinding.yaml” command.

```
Secret.yaml
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: admin-user
5   namespace: kubernetes-dashboard
6   annotations:
7     kubernetes.io/service-account.name: "admin-user"
8 type: kubernetes.io/service-account-token
9
```

We can also create a token with the secret which bound the service account and the token will be saved in the Secret after applying the service using “kubectl apply -f secret.yaml” command.

The implementation of above yaml files can be seen in the below figure.

```
student@edge:~$ kubectl apply -f dashboard.yaml
serviceaccount/admin-user unchanged
student@edge:~$ kubectl apply -f ClusterRoleBinding.yaml
clusterrolebinding.rbac.authorization.k8s.io/admin-user unchanged
student@edge:~$ kubectl apply -f secret.yaml
error: the path "secret.yaml" does not exist
student@edge:~$ kubectl apply -f Secret.yaml
secret/admin-user unchanged
```



The screenshot shows a terminal window with a dark background. At the top, there is a menu bar with the options: File, Actions, Edit, View, and Help. Below the menu bar, there are five tabs, each labeled 'student@edge: ~'. The first tab is active. The terminal prompt is 'student@edge:~\$'. The user has entered the command 'kubectl proxy --address=localhost:8080 --port=8080 --kubeconfig=/home/edge/.kube/config --api-prefix=/apis --api-version=v1 --api-group=core --api-resource=services --api-resource=namespaces --api-resource=pods --api-resource=replicasets --api-resource=secrets --api-resource=storageclasses --api-resource=volumeattachments --api-resource=volumeattachments/status --api-resource=volumeattachments/status/attach --api-resource=volumeattachments/status/detach --api-resource=volumeattachments/status/attach/detach --api-resource=volumeattachments/status/attach/detach/status'. The output of the command is 'Starting to serve on 127.0.0.1:8001'. Below this, there are five lines of log output, each starting with 'E1127 13:59:39.541924 20252 proxy_server.go:147] Error while proxying request: context canceled'. The terminal window also shows a sidebar on the left with the following items: 'Workload Status', 'Pods', 'Containers', 'Deployments', 'Jobs', and 'Pods'.

```
File Actions Edit View Help
student@edge: ~ x student@edge: ~ x student@edge: ~ x student@edge: ~ x student@edge: ~ x
student@edge:~$ kubectl proxy --address=localhost:8080 --port=8080 --kubeconfig=/home/edge/.kube/config --api-prefix=/apis --api-version=v1 --api-group=core --api-resource=services --api-resource=namespaces --api-resource=pods --api-resource=replicasets --api-resource=secrets --api-resource=storageclasses --api-resource=volumeattachments --api-resource=volumeattachments/status --api-resource=volumeattachments/status/attach --api-resource=volumeattachments/status/detach --api-resource=volumeattachments/status/attach/detach --api-resource=volumeattachments/status/attach/detach/status
Starting to serve on 127.0.0.1:8001
E1127 13:59:39.541924 20252 proxy_server.go:147] Error while proxying request: context canceled
E1128 14:33:46.700162 20252 proxy_server.go:147] Error while proxying request: context canceled
E1128 14:33:46.710736 20252 proxy_server.go:147] Error while proxying request: context canceled
E1130 10:36:45.509577 20252 proxy_server.go:147] Error while proxying request: context canceled
Workload Status
Pods
Containers
Deployments
Jobs
Pods
```

Edge on LAB000001 - Virtual Machine Connection

File Action Media Clipboard View Help

student@edge: ~


File Actions Edit View Help

student@edge: ~ x student@edge: ~ x student@edge: ~ x student@edge: ~ x

```
student@edge:~$ kubectcl -n kubernetes-dashboard create token admin-user
eyJhbGciOiJIUzU1NiIsImtpZCI6IiBNXXZwUHFVVGZ5TTJaeklid2Z4Tkc3M21TVG95b212Vm9iaHo3LXNRV0U1fQ.eyJhdwQiOi01siaHR0cHM6Ly9vL2t1YmVybWV0ZXMuZGVmYXV5dC5zdmM1LjCjrdWJlcm5ldGVzLmVlYj7Im5hbWVzcGFjZSI6Imt1YmVybWV0ZXMtZGFzaGJvYXJkIiwic2VydmljZWYiOiF0eSImSjZiIjE6MTcwMTA5MzQ3Nywic3ViIjoic3lzdGltOnNlcnZpY2VhY2N2dW500mt1YmVybWV0ZXMtZGFzaGJvYXJkOmFkbWluLXVzZXIiOiF0.qdUPW0sMtIpMwPGPwVludgt8tV4jvz-DobMVZeIjflfQhg o5tCHLYBSs4mQ Dnyafniaiaxry ljaNRAXH2iDYxIY5CB3 HLGvXGYIeXpwqfQkXcj6AurQOLdJFVMU1ofU15czbt3dmqlfBbDnWdg
student@edge:~$
```

CPU Usage

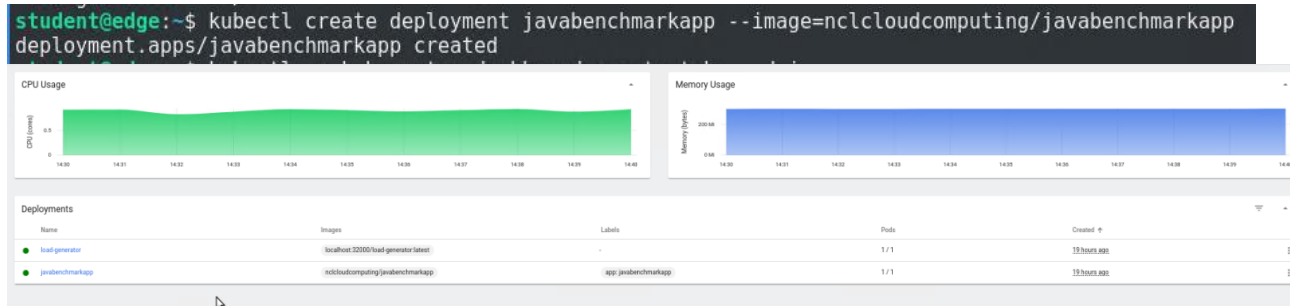
```
kubectl -n Kubernetes-dashboard create token admin-user
```



DEPLOY AN INSTANCE OF THE DOCKER IMAGE "NCLCLOUDCOMPUTING/JAVABENCHMARKAPP" VIA CLI.

We deploy docker image nclcloudcomputing/javabenchmarkapp using the command
 “kubectl create deployment javabenchmarkapp --
 image=nclcloudcomputing/javabenchmarkapp”

Where javabenchmarkapp is the deployment name.

*DEPLOY A NODEPORT SERVICE SO THAT THE WEB APP IS ACCESSIBLE VIA HTTP://LOCALHOST:30000/PRIMECHECK. THE CONTAINER USES PORT 8080 INTERNALLY.*

```

javabenchmarkapp-se...
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: javabenchmarkapp-service
5 spec:
6   type: NodePort
7   selector:
8     app: javabenchmarkapp
9   ports:
10    - protocol: TCP
11      port: 8080
12      targetPort: 8080
13      nodePort: 30000
14
15
16

```

We created a service javabenchmarkapp-service by applying the below given command. In the above file we pass type as NodePort to expose the application service outside the cluster.

```

student@edge:~$ kubectl apply -f javabenchmarkapp-service.yaml
service/javabenchmarkapp-service created
student@edge:~$ kubectl get services

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.152.183.1	<none>	443/TCP	14d
javabenchmarkapp-service	NodePort	10.152.183.110	<none>	8080:30000/TCP	13s

TASK 2: DEPLOY THE MONITORING STACK OF KUBERNETES

Task Objective: Understand and learn how to deploy a monitoring stack of Kubernetes consisting of Prometheus, metrics server, Grafana.

ENABLE OBSERVABILITY SERVICE FROM MICROK8S ADDONS.

Microk8s is a lightweight Kubernetes distributor which is used to get metrics and for monitoring purpose. The below command is used to enable the observability related add Ons for microk8s. The Observability feature includes addons like Grafana, Prometheus.


```
student@edge:~$ microk8s enable observability
Infer repository core for addon observability
Addon core/observability is already enabled
```

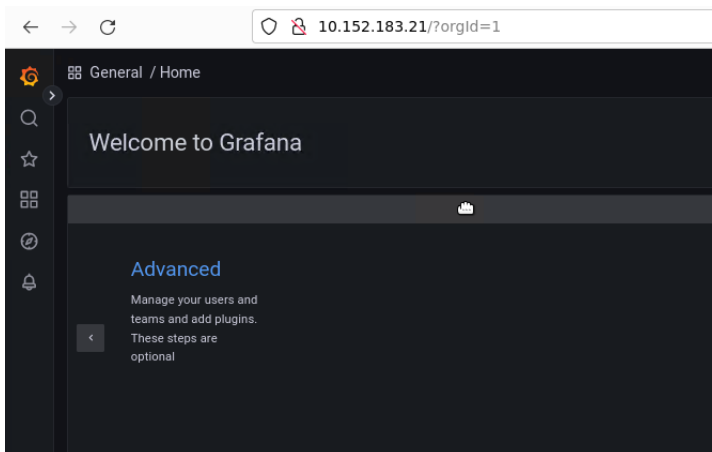
EDIT THE GRAFANA SERVICE TO ALLOW ACCESS FROM THE HOST.

```
student@edge:~$ kubectl edit service/kube-prom-stack-grafana --namespace observability
Edit cancelled, no changes made.
student@edge:~$ kubectl describe service/kube-prom-stack-grafana --namespace observability
Name: kube-prom-stack-grafana
Namespace: observability
Labels: app.kubernetes.io/instance=kube-prom-stack
        app.kubernetes.io/managed-by=Helm
        app.kubernetes.io/name=grafana
        app.kubernetes.io/version=9.3.8
        helm.sh/chart=grafana-6.51.2
Annotations: meta.helm.sh/release-name: kube-prom-stack
             meta.helm.sh/release-namespace: observability
Selector: app.kubernetes.io/instance=kube-prom-stack,app.kubernetes.io/name=grafana
Type: NodePort
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.152.183.21
IPs: 10.152.183.21
Port: http-web 80/TCP
TargetPort: 3000/TCP
NodePort: http-web 31000/TCP
Endpoints: 10.1.11.31:3000
Session Affinity: None
External Traffic Policy: Cluster
Events: <none>
student@edge:~$
```

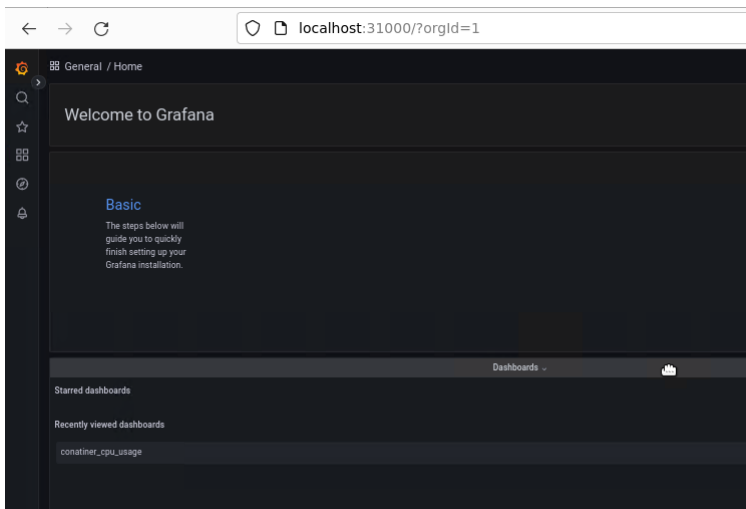
We edit the given Grafana monitoring service, where we change the type of service to NodePort to expose and allow access from the host. We use the command “`kubectl edit service/Kuber-prom-stack-grafana --namespace observability`” to edit the service.

LOG IN TO THE GRAFANA DASHBOARD.

Accessing Grafana service through IP.



Accessing Grafana monitoring service from the host port 31000.



TASK 3: LOAD GENERATOR

Task Objective: Understand the logic of the load generator of benchmarking web applications and the process of deploying your own application of the cluster. Additionally, understand how to build and push a Docker image from scratch.

1. Write a load generator with the following specifications
 - (a) Accepts two configurable values either via a config file or environment variables. target (The address for the load generation) and frequency (Request per second)
 - (b) Generate web request to the target at the specified frequency
 - (c) Collect 2 types of metrics. Average response time and accumulated number of failures
 - (d) Request should timeout if it takes more than 10 seconds. Counted as failures
 - (e) Test results need to be printed to the console
 - (f) There are no requirements in programming language.

```
load_generator.py
1 import requests
2 import time
3 import os
4
5 class LoadGenerator:
6     def __init__(self):
7         self.target = os.environ.get("target", "http://10.152.183.110:8080/primecheck")
8         self.frequency = float(os.environ.get("frequency", 10.0))
9         self.failures = 0
10        self.total_response_time = 0
11        self.total_requests = 0
12
13    def generate_load(self):
14        while True:
15            start_time = time.time()
16            try:
17                response = requests.get(self.target, timeout=10)
18                response.raise_for_status()
19            except requests.exceptions.RequestException as e:
20                print(f"Request failed: {e}")
21                self.failures += 1
22            else:
23                end_time = time.time()
24                response_time = end_time - start_time
25                self.total_response_time += response_time
26                self.total_requests += 1
27                print(f"Request successful. Response Time: {response_time:.2f} seconds")
28
29            time.sleep(1 / self.frequency)
30            self.print_metrics()
31
32    def print_metrics(self):
33        if self.total_requests > 0:
34            avg_response_time = self.total_response_time / self.total_requests
35            print(f"Average Response Time: {avg_response_time:.2f} seconds")
36        print(f"Total Failures: {self.failures}")
37        print(f"Total Requests: {self.total_requests}\n")
38
39 if __name__ == "__main__":
40     load_generator = LoadGenerator()
41     load_generator.generate_load()
42
```

The above load-generator code was created to benchmark our javabenchmarkapp. In the first part of the code, we accept target URL and frequency as input from environment variables, which is further defined in the task4 yaml file. If in case the code is unable to get inputs from our yaml file it would take predefined inputs as defined in the code. In this code we run an infinite loop of sending requests to our javabenchmarkapp to benchmark it. The code sends

invoking request at an interval frequency of 10.0 as defined in the code. We record the response time for each request and also the average response time for all the requests invoked. We use .2f in print statement to restrict how floating point numbers are displayed in response time.

2.AFTER PROGRAMMING, PACK THE PROGRAM AS A STANDALONE DOCKER IMAGE AND PUSH IT TO THE LOCAL REGISTRY AT PORT 32000. NAME THE IMAGE AS LOAD-GENERATOR

```
student@edge:~$ docker run -d -p 32000:3000 --restart=always --name registry registry:2
docker: Error response from daemon: Conflict. The container name "/registry" is already in use by container "eaa05bb131cf3463707bb9a63ed505a464f7e77f0a52bdabd85331285cc365". You have to remove (or rename) that container to be able to reuse that name.
See "docker run --help".
student@edge:~$ docker build -t load-generator .
[*] Building 18.1s (0/9) FINISHED
=> [internal] load .dockerignore
=> transferring context: 2B
=> [internal] load build definition from Dockerfile
=> transferring dockerfile: 103B
=> [internal] load metadata for docker.io/library/python:3
=> [1/4] FROM docker.io/library/python:3@sha256:31ceea009f42df76371a8fb94fa191f988a25847a228dbec35b6f8d2518a6ef
=> [internal] load build context
=> transferring context: 1.58kB
=> CACHED [2/4] WORKDIR /app
=> [3/4] COPY load_generator.py .
=> [4/4] RUN pip install requests
=> exporting to image
=> writing image sha256:40eba16cb58d1dddb7914e52c8384e06e9add884490cdd870b3123e1699481
=> naming to docker.io/library/load-generator
student@edge:~$ docker ps -a
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS                    NAMES
eaa05bb131cf   registry:2     "/entrypoint.sh /etc..." 4 days ago    Up 3 hours    0.0.0.0:32000->5000/tcp, :::32000->5000/tcp   registry
```

Firstly we run docker registry on our host at port 32000 using the docker run command. In the next step we build a docker image of load-generator using the below docker file and running “docker build –t load-generator . “ command.

```
Dockerfile
1 FROM python:3
2
3 WORKDIR /app
4
5 COPY load_generator.py .
6
7 RUN pip install requests
8
9 CMD ["python", "load_generator.py"]
10
```

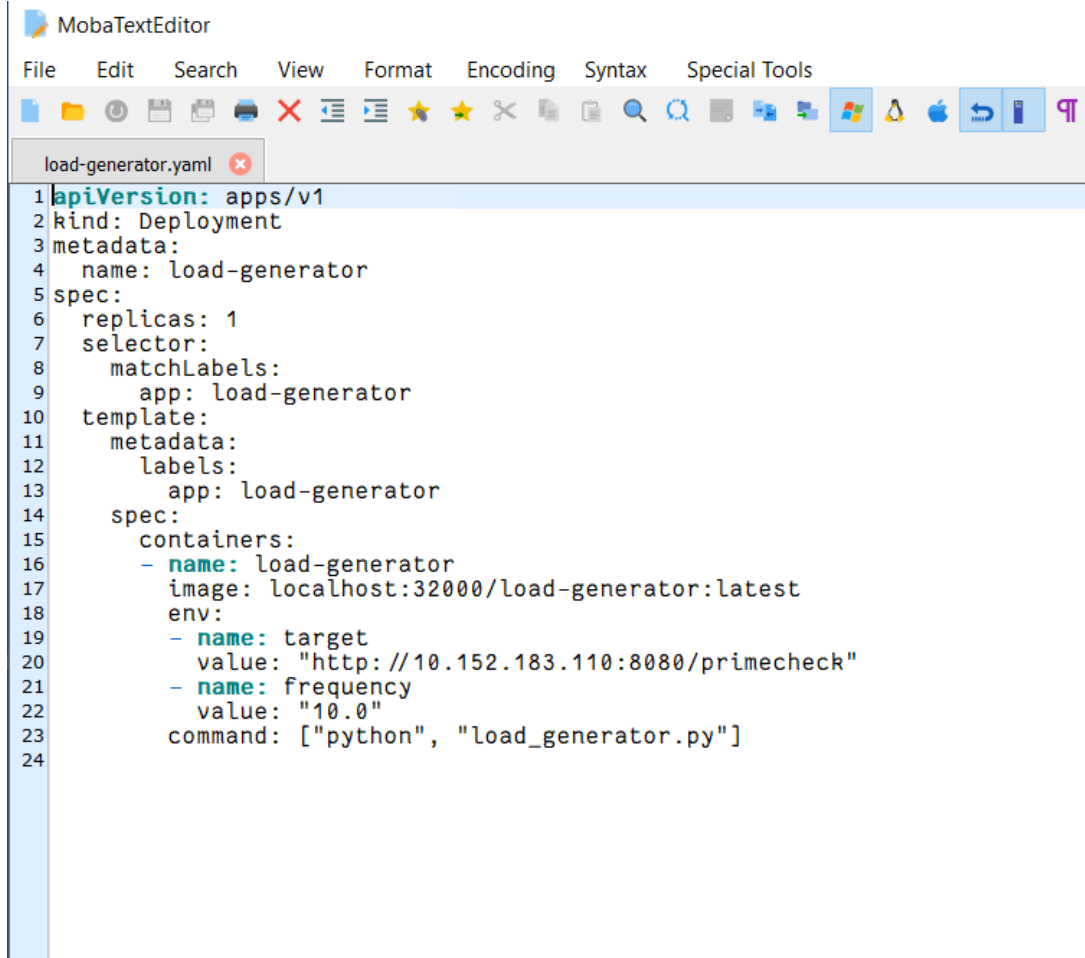
Once we have created the docker image, we tag our image and deploy the same image to our docker registry. To tag the image we use “docker tag load-generator localhost:32000/load-generator“ command. We push the image to docker registry using docker push command.

```
student@edge:~$ docker images
REPOSITORY      TAG        IMAGE ID      CREATED        SIZE
load-generator   latest     40eba16cb58d  38 seconds ago  1.03GB
localhost:32000/load-generator latest     c8c67fcbf020  4 days ago    1.03GB
registry        2         ff1857193a0b  5 weeks ago    25.4MB
student@edge:~$ docker image rm c8c6
Untagged: localhost:32000/load-generator:latest
Untagged: localhost:32000/load-generator@sha256:098aded9a68b43fe142470063d315cefff56584038ae6663f4b2e9b3880e5861
Deleted: sha256:c8c67fcbf0209fc108ea13f591cf3f6de876afd236320e61ea7e744166720941
student@edge:~$ docker tag load-generator localhost:32000/load-generator
student@edge:~$ docker push localhost:32000/load-generator
Using default tag: latest
The push refers to repository [localhost:32000/load-generator]
9b567ab688cd: Pushed
1dc53eb1ad48: Pushed
99aacc8cdf0: Layer already exists
8a90026ddd52: Layer already exists
767b9eb445be: Layer already exists
e47f7ad06b01: Layer already exists
2b4cf8a5bd5e: Layer already exists
80bd043d4663: Layer already exists
30f5cd833236: Layer already exists
7c32e0608151: Layer already exists
7cea17427f83: Layer already exists
latest: digest: sha256:cf280eeee8304344f4f31cb44ad0539c934c9511abb3df8662b0eb8384b75c9 size: 2631
student@edge:~$ kubectl apply -f load-generator.yaml
deployment.apps/load-generator created
```

TASK 4: MONITOR BENCHMARKING RESULTS

Task Objective: To learn and understand how to monitor container metrics.

DEPLOY LOAD-GENERATOR SERVICE CREATED IN TASK 3.



```

MobaTextEditor
File Edit Search View Format Encoding Syntax Special Tools
load-generator.yaml
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: load-generator
5 spec:
6   replicas: 1
7   selector:
8     matchLabels:
9       app: load-generator
10  template:
11    metadata:
12      labels:
13        app: load-generator
14    spec:
15      containers:
16      - name: load-generator
17        image: localhost:32000/load-generator:latest
18        env:
19          - name: target
20            value: "http://10.152.183.110:8080/primecheck"
21          - name: frequency
22            value: "10.0"
23        command: ["python", "load_generator.py"]
24

```

We deploy our load-generator by setting environment variables values that would be used by our load-generator code to continuously send requests to our JavaBenchmarkApp.



```

student@edge:~$ kubectl logs -f load-generator-65dccc8955-hqwxg
Request successful. Response Time: 1.64 seconds
Average Response Time: 1.64 seconds
Total Failures: 0
Total Requests: 1

Request successful. Response Time: 1.70 seconds
Average Response Time: 1.67 seconds
Total Failures: 0
Total Requests: 2

Request successful. Response Time: 1.88 seconds
Average Response Time: 1.74 seconds
Total Failures: 0
Total Requests: 3

Request successful. Response Time: 2.17 seconds
Average Response Time: 1.85 seconds
Total Failures: 0
Total Requests: 4

Request successful. Response Time: 2.74 seconds
Average Response Time: 2.03 seconds
Total Failures: 0
Total Requests: 5

Request successful. Response Time: 2.72 seconds
Average Response Time: 2.14 seconds
Total Failures: 0
Total Requests: 6

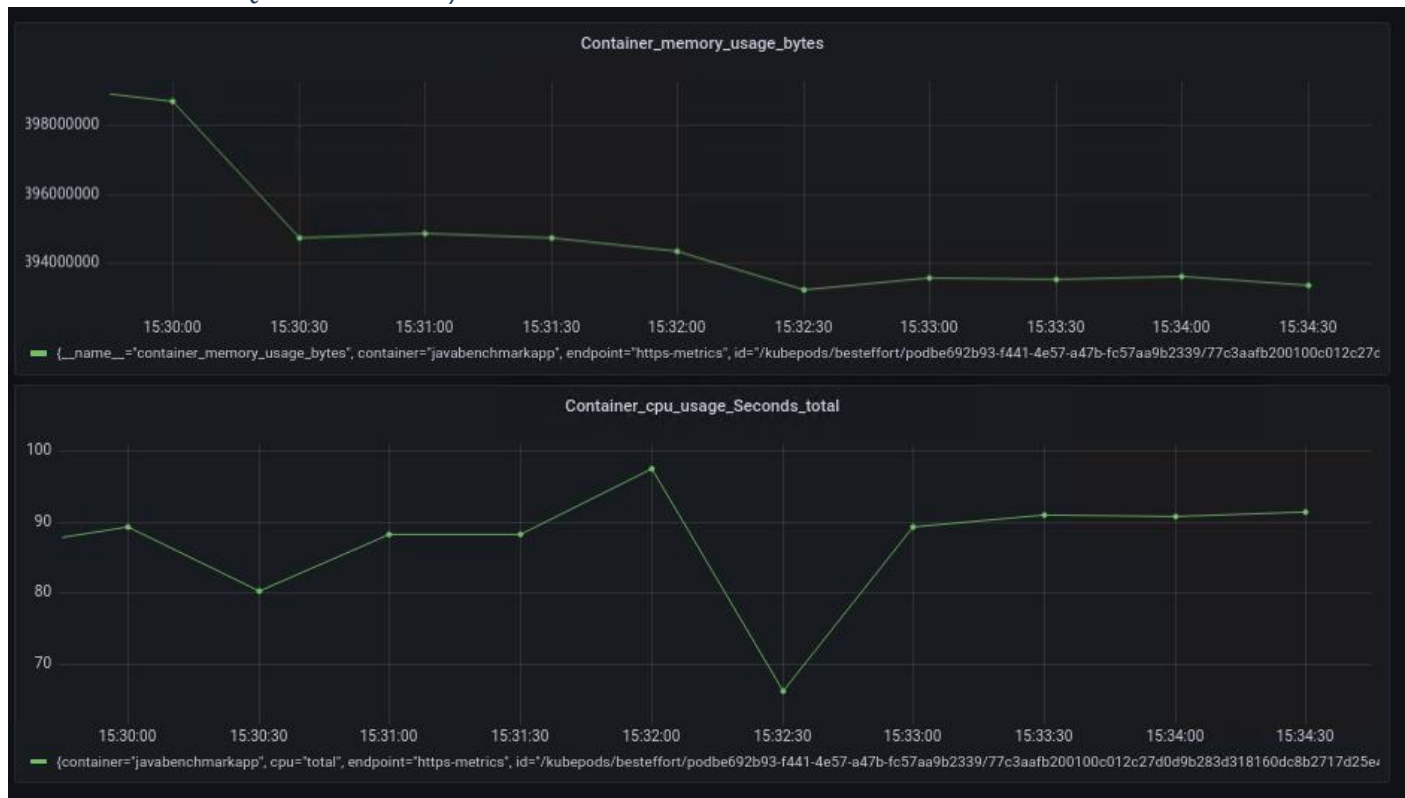
Request successful. Response Time: 2.64 seconds
Average Response Time: 2.21 seconds
Total Failures: 0
Total Requests: 7

Request successful. Response Time: 1.54 seconds
Average Response Time: 2.13 seconds
Total Failures: 0
Total Requests: 8


```

The above figure represents logs obtained after deploying the service.


2. DURING THE BENCHMARKING, CREATE A NEW DASHBOARD ON GRAFANA AND ADD 2 NEW PANELS WHICH SHOULD CONTAIN QUERIES OF CPU/MEMORY USAGE OF THE WEB APPLICATION.



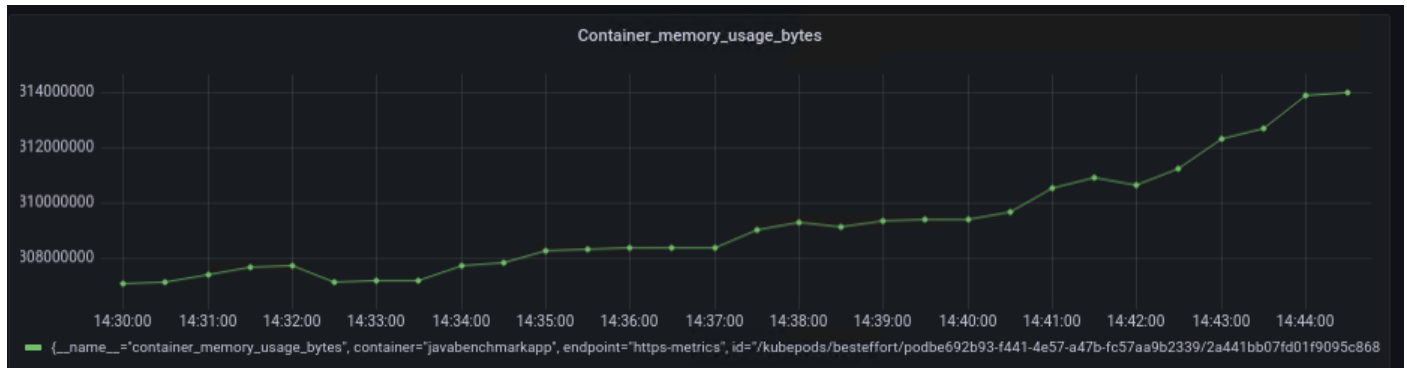
We created a new dashboard and queried the required metrics as shown below.



3. SCREENSHOT THE TWO PANELS



The above image shows the metrics obtained for Container_cpu_usage_Seconds_total, which shows a linear graph that depicts the CPU usage to be increasing with time. The graph shows the usage of CPU once the webapp starts receiving requests, the cpu decline only when the requests fails and it encounters an issue with the service but in our case there is no failure in request, so the graph shows a linear graph.



The above graph shows the Container_memory_usage_bytes metrics obtained which shows non-linear trend of memory usage with time. The Memory usage depends upon the nature of webapp. When the Javabenchmarkapp service is down, the graph shows a negative curve and vice versa when the Javabenchmarkapp is up and running.

CONCLUSION:

Through the above coursework, we were able to obtain practical experience with the deployment, management, and monitoring of Kubernetes applications. We also learned how to generate load for benchmarking, which produces insightful data that can be used to optimize application scaling and enhance overall performance. Using Grafana, we built a dashboard with panels to show the CPU and memory utilization of containers. We discovered that when the load on our webapp increased, the graph displayed a linear curve.

REFERENCES:

1. <https://stackoverflow.com/> : For various doubts related to implementation and errors.
2. <https://kubernetes.io/docs/> : Kubernetes documentation